

Practice 2

Read and Clean

```
readData = function() {  
df = read.csv('titanic.csv', sep=';', header=T)  
df = df[complete.cases(df),]  
df$sex = droplevels(df$sex) # There is a level "" that shouldn't be there.  
df$embarked = droplevels(df$embarked) # There is a level "" that shouldn't be there.  
# http://stackoverflow.com/questions/4605206/drop-data-frame-columns-by-name  
df = subset(df, select = -c(ticket,cabin) )  
# I need to convert all factor features to numbers to produce the correlation matrix.  
# First two functions, applied to each row, substitute the strings by numbers.  
sexToNum = function(x) ifelse((x %in% "male"), 1, 2)  
embToNum = function(x) ifelse(x=="C", 1, ifelse(x=="Q", 2, 3))  
#df$pclass = as.numeric(levels(df$pclass))[df$pclass]  
df$sex = sexToNum(df$sex)  
df$embarked = embToNum(df$embarked)  
df$age = as.numeric(df$age)  
df$fare = as.numeric(df$fare)  
df = df[c("survived", "pclass", "sex", "age", "sibsp", "parch", "fare", "embarked")]  
df  
}
```

Split

```
splitdf = function(dataframe, seed=NULL, percentage=0.8) {  
  if (!is.null(seed)) set.seed(seed)  
  index = 1:nrow(dataframe)  
  numTrainingSamples = round(length(index) * percentage)  
  trainindex = sample(index, numTrainingSamples)  
  trainset = dataframe[trainindex, ]  
  testset = dataframe[-trainindex, ]  
  list(trainset=trainset, testset=testset)  
}
```

Feature Selection

```
easyFeatureSelection = function(split) {  
  corrs = abs(cor(split$trainset)[1,])  
  toKeep = corrs[corrs > 0.1 & !is.na(corrs)]  
  split$trainset = subset(split$trainset, select=names(toKeep))  
  split$testset = subset(split$testset, select=names(toKeep))  
  split  
}
```

Training & Model Evaluation

```
modelEvaluation = function(split, formula=survived~.) {  
  # Fit the model with the training dataset.  
  model = glm(formula, data = split$trainset, family = "binomial")  
  # The predicted probabilities given to each sample in the test set.  
  probs = predict(model, type="response", newdata = split$testset)  
  predictions = data.frame(survived = split$testset$survived, pred=probs)  
  myROC = roc(survived ~ probs, predictions)  
  optimalThreshold = coords(myROC, "best", ret = "threshold")  
  # To compute F1 = 2TP/(2TP+FP+FN)  
  T = table(predictions$survived, predictions$pred > optimalThreshold)  
  F1 = (2*(T[1,1]))/((2*(T[1,1]))+T[2,1]+T[1,2])  
  F1  
}
```

Question 1

Single Splits

```
library(ggplot2)  
library(pROC)  
  
## Type 'citation("pROC")' for a citation.  
##  
## Attaching package: 'pROC'  
  
## The following objects are masked from 'package:stats':  
##  
##      cov, smooth, var  
  
library(caret)  
  
## Loading required package: lattice  
  
df = readData()  
split = splitdf(df, 10, 0.8)  
split = easyFeatureSelection(split)  
singleSplitPerformance = modelEvaluation(split)  
cat("Performance = ", singleSplitPerformance)  
  
## Performance = 0.8434505
```

Iteration

```
perf = c(0.0)  
iterations = data.frame()  
for(i in 1:100) {  
  split = splitdf(df, i, 0.8)  
  split = easyFeatureSelection(split)  
  perf[i] = modelEvaluation(split)  
  row = data.frame(split = i, perf = perf[i])  
}
```


```

iterations = rbind(iterations, row)
}
indexOfMaxPerformance = which.max(perf)
IT_maxPerf = perf[indexOfMaxPerformance]
cat("Max performance = ", IT_maxPerf)

```

```
## Max performance = 0.8806818
```

Comparison: Plot single split vs iteration splits



Compared with the single split, the iterations had a much better performance.

Question 2

Cross validation

Using corss validation, we will divide the data into 60-20-20: 60 for training, 20 for validation and 20 for testing.

On the 20% validation set:

- create a certain **nbfolds** folds
- On each fold: train the model on part of the data and use the remaining part for validation.
- Compute best performance for each fold depending on the chosen method (roc or acc) and store the results in a dataframe

On the 20% testing set:

- Select the best threshold with the best performance from the data frame
- Predict using the same model on the remaining 20% test set and return performance.

The below method will be used to return best performance based on best threshold among **nbfolds** based on both methods: ROC and Accuracy

```
library(ROCR)
```

```
## Loading required package: gplots
```

```

##
## Attaching package: 'gplots'

## The following object is masked from 'package:stats':
##
##      lowess

cross_validation = function(nbfolds, split, method = 'roc'){
  perf = data.frame()
  #create folds
  folds = createFolds(split$trainset$survived, nbfolds, list = TRUE, returnTrain = TRUE)
  #loop nbfolds times to find optimal threshold
  for(i in 1:nbfolds)
  {
    #train the model on part of the data
    model = glm(survived~., data=split$trainset[folds[[i]],], family = "binomial")
    #validate on the remaining part of the data
    probs = predict(model, type="response", newdata = split$trainset[-folds[[i]],])
    if(method == 'roc')
    {
      #Threshold selection based on roc
      #store predictions in data frame
      predictions = data.frame(survived=split$trainset[-folds[[i]],]$survived, pred=probs)
      myROC = roc(survived ~ probs, predictions)
      optimalThreshold = coords(myROC, "best", ret = "threshold",drop=FALSE)
      T = table(predictions$survived, predictions$pred > optimalThreshold[[1]])
      #Measure performace based on best threshold and add performance + threshold to data frame
      F1 = (2*(T[1,1]))/((2*(T[1,1]))+T[2,1]+T[1,2])
      row = data.frame(threshold = optimalThreshold[[1]], accuracy = F1)
    }
    else
    {
      #Threshold selection based on Accuracy instead of ROC
      #create a prediction object based on the predicted values
      pred = prediction(probs,split$trainset[-folds[[i]],]$survived)
      #measure performance of the prediction
      acc.perf = performance(pred, measure = "acc")
      #Find index of most accurate threshold and add threshold in data frame
      ind = which.max( slot(acc.perf, "y.values")[[1]] )
      acc = slot(acc.perf, "y.values")[[1]][ind]
      optimalThreshold = slot(acc.perf, "x.values")[[1]][ind]
      row = data.frame(threshold = optimalThreshold, accuracy = acc)
    }
    #Sote the best thresholds with their performance in the perf dataframe
    perf = rbind(perf, row)
  }

  #Get the threshold with the max accuracy among the nbfolds and predict based on it on the unseen
  indexOfMaxPerformance = which.max(perf$accuracy)
  optThresh = perf$threshold[indexOfMaxPerformance]
  probs = predict(model, type="response", newdata = split$testset)
  predictions = data.frame(survived=split$testset$survived, pred=probs)
  T = table(predictions$survived, predictions$pred > optThresh)
  F1 = (2*(T[1,1]))/((2*(T[1,1]))+T[2,1]+T[1,2])
  F1
}

```

Do 100 iterations:

- Split the data randomly in test and train
- run cross validation on 20% of the data and test on the remaining 20%
- Store the results in a dataframe

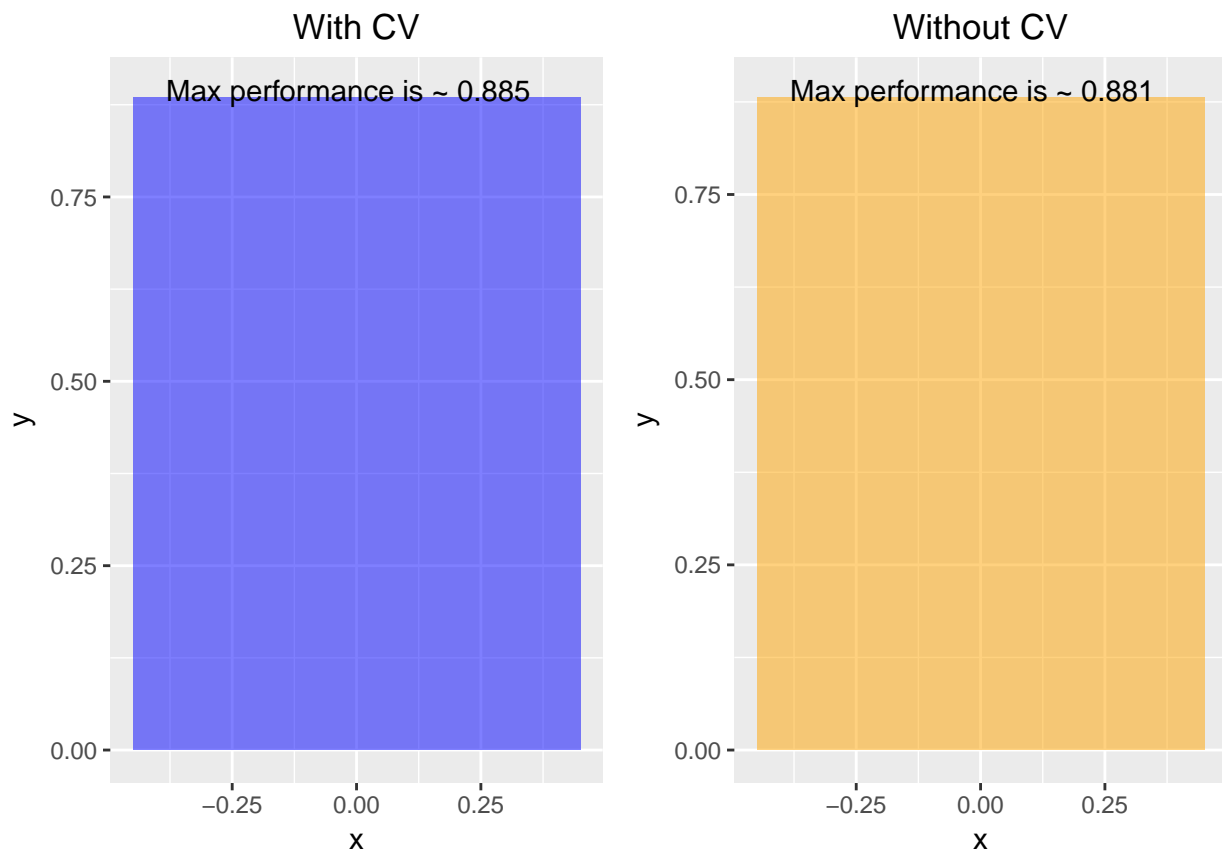
```
#Split the data into 60-20-20, perform cross validation on validation set and evaluate on test set.
sol = data.frame()
data = data.frame()
for(i in 1:100){
  split = splitdf(df, i, 0.6)
  #Choose variables on the 60% training set
  split = easyFeatureSelection(split)
  #divide the remaining 60% into 50% validation (for cv) and 50% for testing
  testvalset = split$testset
  testsplit = splitdf(testvalset, i, 0.5)
  perf = cross_validation(10, testsplit)
  data = data.frame(split = i, performance = perf)
  sol = rbind(sol, data)
}
indexOfMaxPerformance = which.max(sol$performance)
CV_maxPerf = sol$performance[indexOfMaxPerformance]
```

Comparison: with and without CV

We can see that with CV the performance has slightly increased compared to no CV.

```
library(gridExtra)
indexOfMaxPerformance = which.max(sol$performance)
maxPerf = sol$performance[indexOfMaxPerformance]
d1= data.frame(x=c(0), y=c(CV_maxPerf))
d2= data.frame(x=c(0), y=c(IT_maxPerf))
p1= ggplot(data=d1, aes(x, y)) +
  geom_bar(stat="identity", fill = 'blue', alpha = 0.5)+
  annotate("text", label= paste(c("Max performance is ~"), round(CV_maxPerf,3), " "), x= 0, y = CV_maxPerf)
  ggtitle("With CV") +
  theme(plot.title = element_text(hjust = 0.5))
p2 = ggplot(data=d2, aes(x, y)) +
  geom_bar(stat="identity", fill = 'orange', alpha = 0.5) +
  annotate("text", label= paste(c("Max performance is ~"), round(IT_maxPerf,3), " "), x= 0, y = IT_maxPerf)
  ggtitle("Without CV") +
  theme(plot.title = element_text(hjust = 0.5))

grid.arrange(p1, p2, ncol=2)
```



Question 3

Use Accuracy as a measure instead of ROC to choose optimal threshold. Same as above: Do 100 iterations and continuously split the data into train and test sets.

```
sol1 = data.frame()
for(i in 1:100){
  split = splitdf(df, i, 0.6)
  split = easyFeatureSelection(split)
  testvalset = split$testset
  testsplit = splitdf(testvalset, i, 0.5)
  perf1 = cross_validation(10, testsplit, 'acc')
  data2 = data.frame(split = i, performance = perf1)
  sol1 = rbind(sol1, data2)
}
```

Comparison: CV with ROC vs Accuracy as measures

```
indexOfMaxPerformance1 = which.max(sol1$performance)
Acc_maxPerf = sol1$performance[indexOfMaxPerformance1]

p1 = ggplot(data = sol, aes(y=performance, x=split)) +
  xlab("Iteration") +
  ylab("Performance") +
  xlim(0,100) +
  geom_point(alpha = 0.5, color = "blue") +
  geom_hline(yintercept = maxPerf, color = "red") +
```

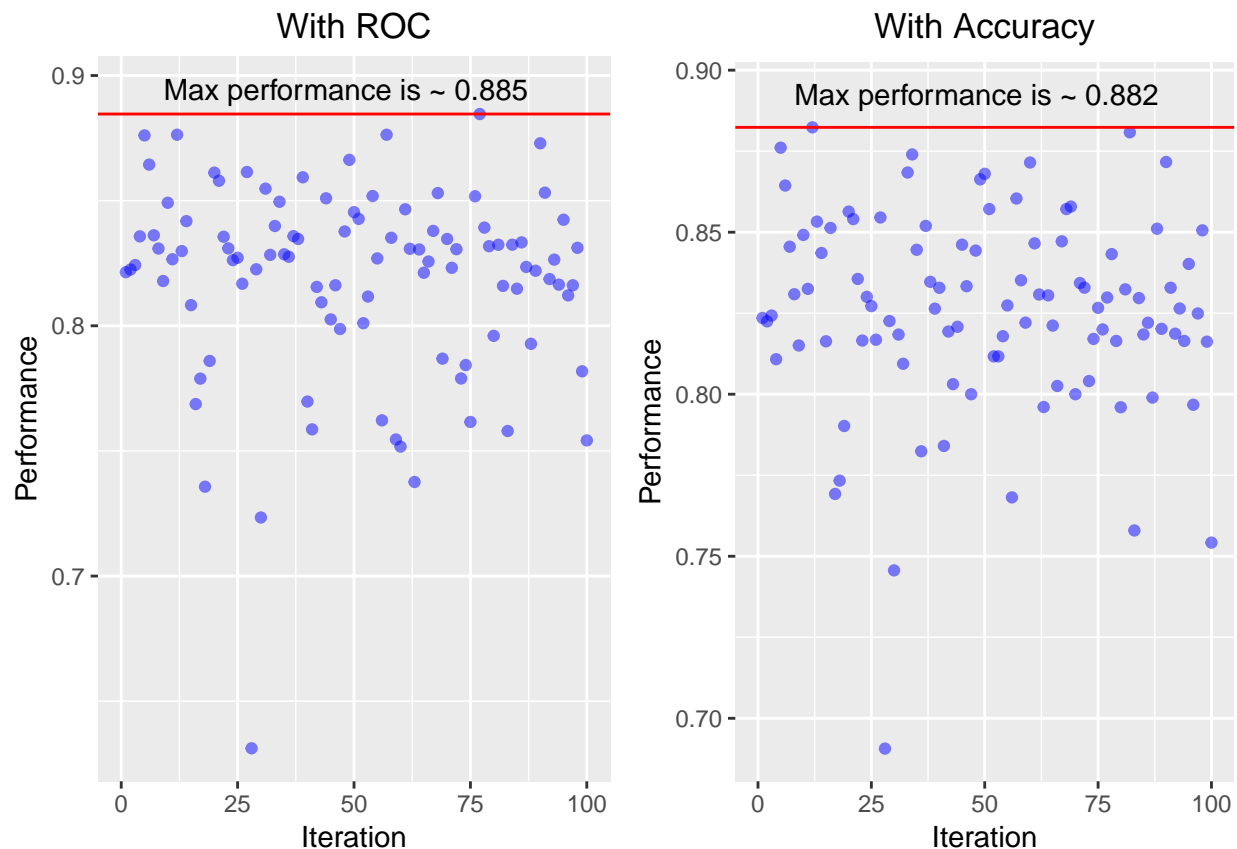
```

    annotate("text", label= paste(c("Max performance is ~"), round(CV_maxPerf,3), " "), x= 50, y = CV_maxPerf) +
    ggtitle("With ROC") +
    theme(plot.title = element_text(hjust = 0.5))

p2 = ggplot(data = sol1, aes(y=performance, x=split)) +
  xlab("Iteration") +
  ylab("Performance") +
  xlim(0,100) +
  geom_point(alpha = 0.5, color = "blue") +
  geom_hline(yintercept = Acc_maxPerf, color = "red") +
  annotate("text", label= paste(c("Max performance is ~"), round(Acc_maxPerf,3), " "), x= 50, y = Acc_maxPerf) +
  ggtitle("With Accuracy") +
  theme(plot.title = element_text(hjust = 0.5))

grid.arrange(p1, p2, ncol=2)

```



The performance of the ROC measure is slightly better than the Accuracy measure.